

intent_router

- Status: PASS

Metrics

Metric	Value
accuracy	0.8077
macro_precision	0.6149
macro_recall	0.6745
macro_f1	0.6114
invalid_label_rate	0.0288
guardrail_accuracy	0.9853
guardrail_precision	0.9709
guardrail_recall	1
guardrail_f1	0.9852
guardrail_false_negative_rate	0
guardrail_false_positive_rate	0.0288

Metric Definitions

Metric	Calculation
accuracy	$\text{intent_accuracy} = \frac{\text{correct_intent_predictions}}{\text{total_intent_samples}}$
macro_precision	$\text{macro_precision} = \text{arithmetic mean of per-class precision across all intent labels}$

Metric	Calculation
macro_recall	macro_recall = arithmetic mean of per-class recall across all intent labels
macro_f1	macro_f1 = arithmetic mean of per-class F1 across all intent labels
invalid_label_rate	invalid_label_rate = invalid_predicted_labels / total_intent_samples
guardrail_accuracy	guardrail_accuracy = (TP + TN) / (TP + TN + FP + FN), with block as positive class
guardrail_precision	guardrail_precision = TP / (TP + FP), with block as positive class
guardrail_recall	guardrail_recall = TP / (TP + FN), with block as positive class
guardrail_f1	guardrail_f1 = 2 * precision * recall / (precision + recall)
guardrail_false_negative_rate	guardrail_false_negative_rate = FN / (TP + FN)
guardrail_false_positive_rate	guardrail_false_positive_rate = FP / (FP + TN)

Guardrail Confusion Matrix

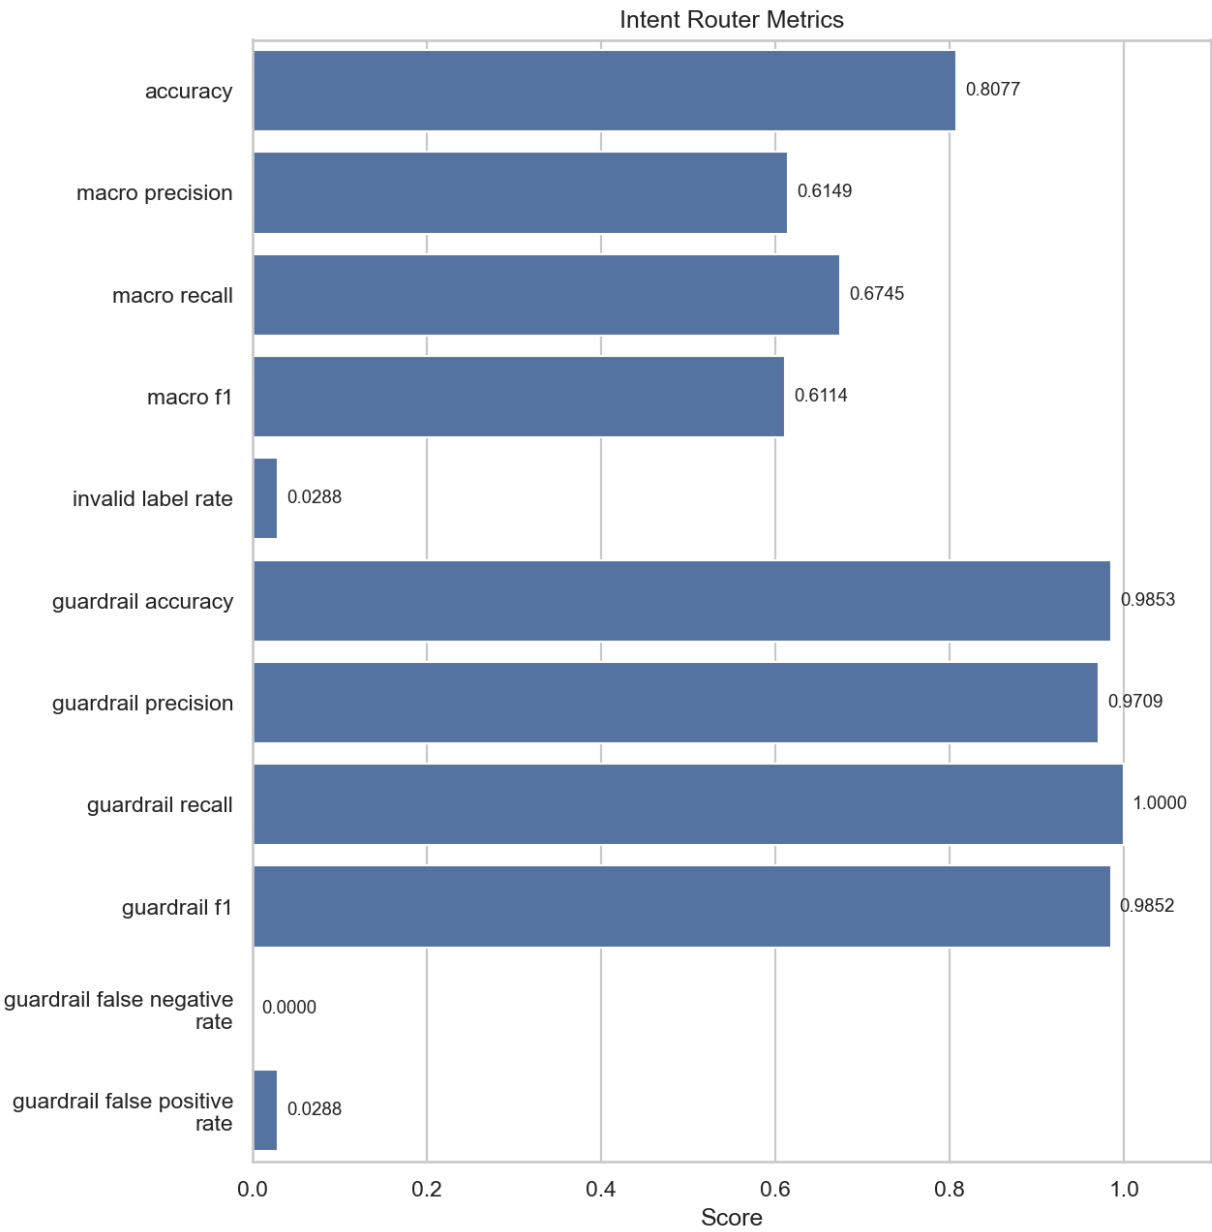
Actual \ Predicted	Allow	Block
Allow	101	3
Block	0	100

Intent Confusion Matrix

Actual \ Predicted	Nursing	Problem Solving	Browsing	Unclear
Nursing	35	1	0	7
Problem Solving	2	26	2	1
Browsing	1	2	18	1

Actual \ Predicted	Nursing	Problem Solving	Browsing	Unclear
Unclear	0	0	0	5

Charts



Notes

- Uses guardrail-first deterministic routing.
- Intent labels are compared against fixture regression data.
- Guardrail metrics use binary labels: `block` vs `allow` .